

Task 4: Training Report

Evaluation Methodology

For this task, both the original wav2vec2 model (from Task 2a) and the fine-tuned model (from Task 3) were evaluated on the cv-valid-dev dataset. The evaluation used Word Error Rate (WER) as the primary metric, where accuracy is calculated as:

$$\text{Accuracy} = 1 - \text{WER}$$

This metric was computed for each audio sample transcription, providing a quantitative comparison of the models' performance.

Results & Observations

- **Original Model:** Average accuracy = **0.883**
- **Fine-tuned Model:** Average accuracy = **0.883**

While the average accuracy for both models is identical, a deeper inspection of the transcriptions reveals differences in performance across individual samples. The fine-tuned model demonstrated improvement on certain clips. However, it also underperformed on other samples, compared to the original model.

Explanation: Overfitting Due to Small Training Set

I hypothesize that this discrepancy can be attributed to overfitting. Due to time constraints, I was limited to training the fine-tuned model on a relatively small subset of the full dataset. Specifically, I trained the model using only 350 training samples over 2 epochs.

Small Training Set (350 samples)

Fine-tuning a large, pre-trained model like wav2vec2 with a limited amount of data (350 samples) can lead to overfitting. The model may become too specialized in the patterns found in the small training set, leading to memorization of specific audio characteristics. As a result, the model performs well on the training set but struggles to generalize to new, unseen examples—like the samples from the cv-valid-dev dataset.

Small Validation Set (150 samples)

A limited validation set (150 samples) compounds the issue. If the validation data lacks sufficient diversity in terms of speech patterns, accents, and background noise, the model can be misled into thinking it has learned all the necessary features. The

validation set may not represent the full variability of real-world data, making it harder for the model to generalize beyond the narrow patterns it has seen.

Proposed Solutions to Improve Accuracy

To further improve the accuracy of the fine-tuned model, the following steps are recommended:

1. Increase Training & Validation Data Size:

Train the model on a larger dataset. Ideally, use the full cv-train-dev dataset or increase the training set with additional data from other sources, such as: LibriSpeech, TED-LIUM, VoxPopuli

Ensure that the validation set is larger and more representative of the target distribution. It should cover a variety of speakers, accents, and environmental conditions to provide a better estimate of model performance in real-world scenarios.

2. Data Augmentation:

Implement data augmentation techniques to artificially expand the training set, including:

- **Noise addition** (background noise, reverberation)
- **Pitch and speed variations**
- **Volume shifts** These transformations would force the model to learn more robust features, helping it perform better on unseen data.

3. Training Hyperparameters:

Standardize preprocessing across both training and evaluation data:

- **Batch sizes** currently due to computational limitations I am running a batch size of 2. A smaller batch size can introduce regularization effects due to noisy updates, which might help in low-data scenarios like mine. A larger batch size might stabilises the gradient updates but may require more data to avoid overfitting.
- **Number of Epochs**, More epochs allow the model to learn more, but with small datasets, this often leads to overfitting

4. Hyperparameter Search: Using Random Search:

Random search is better than a grid search in most cases. Instead of trying every possible combination like grid search, random search samples random combinations. It is more efficient as shown from studies like Bergstra & Bengio, 2012

Use random search to explore different combinations of hyperparameters. This involves defining a set of possible values for each hyperparameter and randomly

sample combinations of these values. Train and evaluate a model for each sampled combination. For example, a random search might explore variations of:

- **Learning rate:** [5e-5, 3e-5, 1e-5]
- **Batch size:** [8, 16, 32]
- **Dropout rate:** [0.1, 0.2, 0.3]
- **Number of epochs:** [2, 3, 4]

Benefits:

- Helps identify optimal training settings, especially when model performance is sensitive to changes.
- Can reveal interactions between parameters (e.g., a higher learning rate working well only with a smaller batch size).
- Provides empirical evidence for hyperparameter selection, rather than relying on defaults or trial-and-error.

6. Evaluate with Additional Metrics:

- Supplement accuracy and WER with other evaluation metrics:
 - **Character Error Rate (CER)** for finer-grained performance analysis.
 - **Semantic similarity metrics** like BLEU or ROUGE to evaluate how well the model captures meaning in transcriptions.

Evaluated both the original model and the fine-tuned model on the cv-valid-dev dataset.

I used Word Error Rate (WER) to determine the accuracy. (Accuracy = $1 - \text{WER}$) This was computed for each audio sample transcription.

Results & observation:

Original model average accuracy = 0.883

Fine-tuned model average accuracy = 0.883

Upon inspecting the transcriptions, both models produced different results. The fine-tuned model improved on some clips but did worse on others compared to the original model. I assume the fine-tuned model is overfitting to the data, as unfortunately due to time constraints I opted to train the fine-tuned model on a relatively small subset of the full dataset. I trained it on 350 samples over 2 epochs. This could likely have cause

Explanation:

Overfitting Explanation:

Small Training Set (350 samples):

Fine-tuning a large model like wav2vec2 with a small training set can cause the model to memorize the training data rather than generalize well to unseen data. This leads to overfitting, where the model performs well on the training set but struggles with new, unseen examples (like the cv-valid-dev dataset).

Validation Set (150 samples):

A small validation set further compounds this issue. If the validation set doesn't cover enough diversity in speech patterns, accents, and audio conditions, it can mislead the model into adapting only to the narrow patterns seen in the validation set, rather than learning broader generalizable features.

Results of Overfitting:

The fine-tuned model may perform well on some examples (similar to what it saw during training) but worse on others that differ slightly from the training data (e.g., different speakers, accents, or noise levels).

The discrepancy in performance (better on some, worse on others) is typical of overfitting, where the model has effectively memorized parts of the training data but fails to generalize.

Solution to increase accuracy:

Further improvement to accuracy:

(other models or datasets/ methods)

- This suggests the fine-tuning process introduced adaptations specific to the training subset, without

significantly generalizing.

Observations:

- Despite the same average accuracy, each model excelled in different scenarios.

- Fine-tuning with limited data may have led to overfitting.

- Trade-offs between samples balanced the accuracy score.

Steps to Improve Accuracy:

1. Use the full Common Voice training set for fine-tuning.

2. Include additional datasets such as LibriSpeech, TED-LIUM, and VoxPopuli.

3. Apply data augmentation: background noise, pitch/speed shift, and volume changes.

4. Normalize text preprocessing across all stages (e.g., case folding, punctuation handling).

5. Experiment with model fine-tuning strategies: learning rate schedules, early stopping, partial layer freezing.

6. Evaluate using other metrics like Character Error Rate (CER), BLEU, and ROUGE.

7. Use active learning to target and collect samples where WER is high.

Conclusion:

The fine-tuned model showed potential for domain-specific improvements but did not outperform the base

model overall. Strategic data scaling, augmentation, and optimization are needed for further gains.